

Towards Social Profile Based Overlays

David Isaac Wolinsky, Pierre St. Juste, P. Oscar Boykin, Renato Figueiredo
University of Florida

Abstract

Online social networking has quickly become one of the most common activities of Internet users. As social networks evolve, they encourage users to share more information, requiring the users, in turn, to place more trust into the social network. Peer-to-peer (P2P) overlays provide an environment that can return ownership of information, trust, and control to the users, away from centralized third-party social networks.

In this paper, we present a novel concept, social profile overlays, which enable users to share their profile only with trusted peers in a scalable, reliable, and private manner. Each user's profile consists of a unique private, secure overlay, where members of that overlay have a friendship with the overlay owner. Profile data is made available regardless of the state of the profile owner through the use of the profile overlay's distributed data store. Privacy and security are enforced through the use of a public key infrastructure (PKI), where the role of certificate authority (CA) is handled by the overlay owner and each member of the overlay has a CA-signed certificate. To discover friends and bootstrap connections into the profile overlay, each member of the social network joins a common public or directory overlay. We define interfaces and present tools that can be used to implement this system, as well as explore some of the challenges related to it.

1 Introduction

Online social networking has become pervasive in daily life, though as social networks grow so does the wealth of personal information that they store. Once information has been released on a social network, known as a user's profile, the data and the user are at the mercy of the terms dictated by the social network infrastructure, which today is typically third-party, centrally owned. If the social network engages in activities disagreeable to the user, due to change of terms or opt-out programs not well understood by users such as recent issues with Facebook's Beacon program [16], the options presented to the user are limited: to leave the social network (surrendering their identity and features provided by the social network), to accept the disagreeable activities, or to petition and hope that the social network changes its behavior.

As the use of social networking expands to become the primary way in which users communicate and express their identity amongst their peers, the users become more dependent on the policies of social network infras-

tructure owners. Recent work [3] explores the coupling between social networks and P2P systems as a means to return ownership to the users, noting that a social network made up of social links is inherently a P2P system with the aside that they are currently developed on top of centralized systems. In this paper, we extend this idea with focus on the topic of topology; that is, how to self-organize social profiles that leverage the benefits offered by a structured P2P overlay abstraction.

Structured P2P overlays provide a scalable, resilient, and self-managing platform for distributed applications. Structured overlays enable users to easily create their own decentralized systems for the purpose of data sharing, interactive activities, and other networking-enabled activities. In recent work [24], we have implemented mechanisms that allow users to create and manage their own private overlays using a common public overlay to assist in discovery and NAT traversal. This prior work focuses on generic structured P2P private overlays; in this paper, we expand upon this approach with in-depth discussion on how to apply this technique to enable social network overlay profiles.

Social networks consist of users, each of whom has a profile, a set of friends, and private messaging. The profile contains the user's personal information, status updates, and public conversations. Friends are individuals which the user trust sufficiently to view their profile. Private messaging allows users to send messages between each other without leaking any information to other friends. Using this model, we describe how a common directory overlay can be used to provide services for finding friends and joining existing profile overlays. Each user has their own profile overlay, secured via public key infrastructure (PKI) to which they are the certificate authority (CA). The profile data is stored in the profile overlay's distributed datastore, allowing profile information to be accessed in scalable mechanisms regardless of the profile owner's online state. In this paper, we present the architecture of these overlays as presented in Figure 1 and how they are used to find and befriend peers, and describe approaches to handling profile data, and establishing initial connections to profile overlays.

The rest of this paper is organized as follows. Section 2 provides background and related work. Section 3 describes our multi-overlay approach, explaining how to map social networks onto structured P2P overlays. In Section 4, we explore some of the remaining challenges confronted by our system. We conclude the paper in Section 5.




Figure 1: An example social overlay network. Alice has a friendship with Bob and Carol, and thus both are members of her profile overlay. Bob has a friendship with Alice and Dave but not Carol; thus Alice and Dave are members of his profile overlay, while Carol is not. Each peer has many overlay memberships but a single root; this relationship is represented by dashed lines in various shades of grey. For clarity, overlay shortcut connections are not shown.

2 Background

In this section, we review structured P2P overlays and distributed, decentralized online social network approaches.

2.1 Structured P2P Overlays

Structured P2P systems provide distributed lookup services with guaranteed search time with a lower bound of $O(\log N)$, in contrast to unstructured systems, which rely on global knowledge/broadcasts, or stochastic techniques such as random walks [6]. Some examples of structured systems can be found in [20, 22, 14, 15, 17, 10]. In general, structured systems are able to make these guarantees by self-organizing a structured topology, such as a 2D ring or a hypercube.

In the overlay, each node is given a unique node ID drawn from a large address space. Each node ID must be unique otherwise address collisions will occur, which can prevent nodes from participating in the overlay. Furthermore, having the node IDs well distributed assists in providing better scalability as many shortcut selection algorithms depend on having node IDs uniformly distributed across the entire address space. Two approaches to ensure this behavior are to have each node use a cryptographically strong random number generator to generate the node ID, or to use a trusted third party generate node IDs and cryptographically sign them [7].

Overlay shortcuts enable efficient routing in ring-structured P2P systems. Different shortcut selection methods include: maintaining large tables without using connections and only verifying usability when routing messages [20, 15], maintaining a connection with a peer every set distance in the P2P address space [22], or using locations drawn from a harmonic distribution in the

node address space [14].

Most structured P2P overlays support decentralized storage/lookup of information by mapping keys to specific node IDs in an overlay. At a minimum, the data is stored at the node ID either smaller or larger to the data's node ID and for fault tolerance the data can be stored at other nodes. This sort of mapping and data storage is called a distributed hash table (DHT). DHTs provide the building blocks to form more complex distributed data stores as presented in Past [19] and Kosha [5].

In [8, 18], the authors discuss the concept a single overlay supporting services through the use of additional overlays that use the underlying overlay to assist in discovery. In [24], we describe a reference implementation of a multiple overlay system that supports the use of a public overlay's DHT to store currently active peers in the private overlays. The system allows users to create their own private overlays without having to create their own bootstrap network. During evaluation, using both simulated and real systems, we saw that the time for one node to join a public and thereafter private overlay was small and grew logarithmically with network sizes. The real system was tested using a public overlay of 600 nodes and any where from 1 to 600 peers in the private overlay using security, the time to connect was always less than 22 seconds. In simulations, we tested public and private overlays varying from 1,000 to 100,000 peers. For 1,000 public overlay peers, regardless of the private overlay size though with security enabled, connected within 22 seconds. The 100,000 peer system, regardless of the private overlay size though with security enabled, connected in less than 48 seconds.

In addition, our system provides both relay-based and hole-punching NAT traversal [23] techniques and supports point-to-point PKI based security [24], forming a basis for the approach presented in this paper.

2.2 Peer-to-Peer Social Networks

The recent popularity and growing privacy issues of centralized online social networks has motivated research projects aimed at providing private, P2P social networks [4, 9, 1, 21].

In [4], a DHT provides the lookup service for storing metadata pertaining to a peer's profile. Peers query the DHT for updated content from their friends by hashing their unique identifiers (e.g. friends' email addresses). The retrieved metadata contains information for obtaining the profile data such as IP address and file version. Their work relies on a PKI system that provides identification, encryption, and access control. In contrast, our approach provides each user their own private overlay secured by point-to-point encryption and authentication amongst all peers in the profile overlay. The profile overlay provides a clean abstraction of access control, whereby once admitted to a private overlay, users can ac-

cess a distributed data store which holds the contents of the owners profile.

[21] takes a different approach by depending on virtual individual servers (VIS) hosted on a cloud infrastructure such as Amazon EC2. Friends contact each other's VIS directly for updates and uses a DHT as a directory for groups or interest-based searches. Their approach assumes bidirectional end-to-end connectivity between each VIS, where a profile is only available during the uptime of the VIS. Because of the demands on network connectivity and uptime, the approach assumes a cloud-hosted VIS and has difficulty being used on user-owned resources. Our approach enables users to avoid the need for all-to-all connectivity and constant uptime through the use of NAT traversal support and the ability to store the profile in the overlay's distributed data store.

The matryoshka approach presented in [9] also uses a DHT for looking up a peer's matryoshka or circle of friends. Once a node in the peer's outermost circle is found, that node is used to route profile requests to the innermost circle which contains replicas of a peer's profile. Trust is enabled through the use of an identification service contacted through the DHT. The circle of friends concept lacks the simplicity of the abstraction made in our approach, whereby a variety of existing structured overlay techniques can be used as a profile overlay without concern over innermost and outermost circles. Our approach also enables the profile owner to serve as a CA, ensuring that nodes can only access a profile overlay after having obtained a signed certificate.

Unlike the above approaches, the P2P social network presented in [1] uses an unstructured overlay and does not require a DHT. Peers connect to each other directly over IP without any overlay routing. Once peers are friends, they maintain unique identifiers to deal with dynamic IPs. Peers cache each other's data for availability through replication, and helper nodes are used to assist with communication between peers behind NATs. However, the approach lacks security and access control considerations and lacks the guarantees and the simplicity of the abstraction offered by a structured overlay.

3 Social Overlays

In this section, we explain how to map online social networking to our multi-overlay social network consisting of a public directory overlay with many private profile overlays. The directory overlay supports friend discovery and verification and stores a lists of peers currently active in each profile overlay. Profile overlays support message boards, private messages, and media sharing.

3.1 Finding and Verifying Friends

In a traditional social network, a directory provides users the ability to search for other users using public

information, such as the user's full name, user ID, e-mail address, group affiliations, and friends. The search results return zero or more matching directory entries. Based upon the results, the user can potentially make a friendship request. The request receiver can review the public information of the requestor prior to making a decision. If the receiver accepts the request, the peers are given access to each other's profiles. Once profile access has been enabled, the users can learn more information, and if it turns out to be a mistake, the peers can unilaterally end the relationship.

To map this to our proposed social overlay, the directory entries would be inserted into the DHT of a public overlay. As discussed in previous work, the keys where the directory entries are stored at consist of a subset of the user's public information in lower-case format and hashed to an overlay address. The value stored at these keys is the user's certificate, which consists of its public information and an overlay address where the user expects to receive notifications. The overlay address can be used for asynchronous offline messaging, whose function we will explain shortly.

Because the users need a way to verify each other that involves social credentials, we propose the use of a new form of certificate. The main portion of the certificate is similar to a self-signed x509 [13] certificate with public information such as user's name, e-mail address, and group affiliations embeded into the certificate. At the tail of the certificate is a friend list represented by many friend entries. To do this we propose employing a technique similar to PGP [11]: users can acquire from their friends a signed message consisting of a hash of the peer's certificate, the time stamp, and the friend's certificate hash signed by the friend. Since PGP does not provide a strong method for revocation, the time stamp provides additional information to help decide whether or not a friendship link is still active without accessing the profile overlay of either peers. Peers should request a new friend list entry within a certain period of time or it will appear that the friendship is no longer valid.

While looking for an individual, a peer may discover that many individuals have overlapping public information components, such as the user's name. Assuming all entries are legitimate, the overlay must support inserting multiple values at the same key, leaving the peer or the peer's DHT client to parse the responses and determining the best match by reviewing the contents of each certificate. Alternatively, a technique like Sword [2], which supports distributing the data across a set of nodes and using a bounded broadcast to discover peers that match all information, could be used for searching.

Upon discovering an individual with whom a peer would like to establish a friendship, the peer will issue a friendship request. As stated earlier, the data stored

in the directory has an overlay address, where a peer expects friendship requests to be inserted into the DHT. The friendship request consists of the self-signed certificate of the requesting peer, the public information of the request receiver, and a time stamp, all signed with the private key associated with the requesting peer's private key matched to their self-signed certificate.

Within a reasonable amount of time after a request has been inserted into the DHT, the receiver can come online and check for outstanding requests. If the receiver would like to add the user, they may do so conditionally or unconditionally. During an unconditional accept, the peer signs the request of the originating requester and issues a request to befriend the originating requester. Alternatively in the case of a conditional accept, the user issues a friendship request, waits for a reply, and investigates the profile prior to signing the originating requester's request. Once a user has received a signed certificate, they may access the remote peer's profile overlay as discussed in 3.2, which is also responsible for activities such as revocation.

3.2 The Profile Overlay

In a traditional social network, the profile or user-centric portion consists of a public message board for status updates or public messages, private messaging, data sharing, and maintenance of existing friendships. In this section, we explain how these components can be applied to a structured overlay dedicated to an individual profile.

The profile overlay consists of all the online friends of the profiler owner. Using the techniques such as those described in [24], it is feasible to efficiently multiplex a P2P across multiple, virtual private overlays. For access control, we employ a PKI, where each member uses the signed certificate generated during the "finding and verifying friends" stage. All links are encrypted using symmetric security algorithms established through the PKI, thus preventing uninvited guests from gaining direct access to the overlay and hence the profile. Because the profile owner also is the CA for all members of the overlay, they can easily revoke users from access to the profile overlay. In [24], we described mechanisms for overlay revocation through the use of broadcasting for immediate revocation and the use of DHT for indirect and permanent revocation.

The message board of a profile can be stored in two ways: distributed within the profile overlay via a data store or stored on the profile owner's personal computing devices. The distributed data store provides the profile when the owner is offline and also distributes the load for popular profiles. For higher availability, each peer should always be a provider for all data in their profile when they are online. To ensure authenticity and integrity, all peers should sign their messages and each peer's certificate should be available in the overlay for

verification. Messages that are unsigned should be ignored by all members of the overlay. An ideal overlay for this purpose should support complex queries [12] allowing easy access to data stored chronologically, by content, by byte type, i.e., media, status updates, or message board discussions.

Private messaging in the profile overlay is unidirectional meaning that only the profile owner can receive private messages using their overlay. To enforce this, a private message should be prepended with a symmetric key encrypted by the profile owner's public key, the message should be appended by a hash of the message to ensure integrity and the entire message encrypted by the symmetric key. This approach ensures that only the sender and the profile owner can decrypt the private message. The contents of the private message should include the sender, time sent, and the subject. Messages can be stored in well known locations, so that the profile owner can either poll the location or, alternatively, use an event based system to notify them of the new message.

3.3 Active Peers

The directory overlay should be used to assist in finding currently active peers in the profile overlays. By placing their node IDs at a well-known, unique per-profile overlay key in the DHT, active peers can bootstrap incoming peers into the profile overlay. We implemented and evaluated this concept in [24]. Because the profile overlay members all use PKI to ensure membership, even if malicious peers insert their ID into the active list, it would be useless as the peer would only form connections with peers who also have a signed certificate for that profile.

4 Challenges

The following list consists of the issues that need further addressing, though list is not exhaustive:

1) Handling small overlay networks - Most research into overlay focuses on networks larger than the typical user's friend count, the average on Facebook is 125. This comes up with regards to the reliability of the overlay as the system may potentially be under constant churn. This is why it is critical that a user host their own profile, but what happens if the user is disconnected? Churn is so significant that a profile cannot be maintained in the overlay? Perhaps a group of users could come together using a common, centralized resource to provide dedicated hosting for their social profiles. This keeps the spirit of user owned systems and contributes to the reliability of the social overlay.

2) Overlay support for low bandwidth, unconnected devices - devices such as smart phones cannot constantly be actively connected to the overlay and the connection time necessary to retrieve something like a

phone number may be too much to make this approach unuseful. Similar to the previous challenge, this approach could benefit from having shared common, centralized resource enabling users access to their social overlays by proxy without establishing a direct connection to the overlay network.

3) Reliability of the directory and profile overlay

- All overlays are susceptible to attacks that can nullify the usefulness of the overlay. While the profile overlay does have point-to-point security it can be very difficult to reclaim an overlay if there are enough malicious users. In a public overlay without some form of centralization, policing the system becomes a very complicated procedure. Previous work has proposed methods to ensure the usability of overlays even while under attack. For the social overlay to be successful, we must identify which methods should be used.

5 Conclusion

In this paper, we proposed methods by which a social network can be decentralized through the use of structure P2P overlays. Our system involves the use of multiple overlays where all users join a common public overlay and individual profile overlays. The public overlay provides directory services that enable users to find and befriend other peers and bootstrap connections into the secure profile overlays. Upon forming a friendship through the directory overlay, peers are given CA signed certificates that allow them to join each other's profile overlay. The owner of the profile overlay acts as CA allowing the user to unilaterally revoke certificate, thus ending friendships with members of their overlay using efficient and reliable methods. For the purpose of storing profile information into the overlay, we cite previous work that can be used to provide distributed data services and give examples of how to store data securely in the overlay. Our proposed system returns control of the social network and more importantly users' identity to the users and eliminates the need for centralized social networks.

References

- [1] S. M. A. Abbas, J. A. Pouwelse, D. H. J. Epema, and H. J. Sips. A gossip-based distributed social networking system. In *Enabling Technologies, IEEE International Workshops on*, 2009.
- [2] J. Albrecht, D. Oppenheimer, A. Vahdat, and D. A. Patterson. Design and implementation trade-offs for wide-area resource discovery. In *ACM Trans. Internet Technol.*, 2008.
- [3] S. Buchegger and A. Datta. A case for P2P infrastructure for social networks - opportunities & challenges. In *WONS '09: The Sixth International Conference on Wireless On-demand Network Systems and Services*, 2009.
- [4] S. Buchegger, D. Schiöberg, L. H. Vu, and A. Datta. PeerSON: P2p social networking: early experiences and insights. In *SNS '09: Proceedings of the Second ACM EuroSys Workshop on Social Network Systems*, 2009.
- [5] A. R. Butt, T. A. Johnson, Y. Zheng, and Y. C. Hu. Kosha: A peer-to-peer enhancement for the network file system. In *SC '04: Proceedings of the 2004 ACM/IEEE conference on Supercomputing*.
- [6] M. Castro, M. Costa, and A. Rowstron. Debunking some myths about structured and unstructured overlays. In *NSDI'05: Proceedings of Symposium on Networked Systems Design & Implementation*, 2005.
- [7] M. Castro, P. Druschel, A. Ganesh, A. Rowstron, and D. S. Wallach. Security for structured peer-to-peer overlay networks. In *Symposium on Operating Systems Design and Implementation (OSDI'02)*, December 2002.
- [8] M. Castro, P. Druschel, A.-M. Kermarrec, and A. Rowstron. One ring to rule them all: Service discover and binding in structured peer-to-peer overlay networks. In *SIGOPS European Workshop*, Sept. 2002.
- [9] L. A. Cuttello, R. Molva, and T. Strufe. Privacy preserving social networking through decentralization. In *Wireless On-Demand Network Systems and Services (WONS'09)*, 2009.
- [10] G. DeCandia, D. Hastorun, M. Jampani, G. Kakulapati, A. Lakshman, A. Pilchin, S. Sivasubramanian, P. Voshall, and W. Vogels. Dynamo: amazon's highly available key-value store. In *SOSP '07: Proceedings of twenty-first ACM SIGOPS symposium on Operating systems principles*, pages 205–220, New York, NY, USA, 2007. ACM.
- [11] S. Garfinkel. *PGP: Pretty Good Privacy*. O'Reilly & Associates, Inc., Sebastopol, CA, USA, 1996.
- [12] M. Harren, J. M. Hellerstein, R. Huebsch, B. T. Loo, S. Shenker, and I. Stoica. Complex queries in dht-based peer-to-peer networks. In *IPTPS '01: Revised Papers from the First International Workshop on Peer-to-Peer Systems*, 2002.
- [13] R. Housley, W. Polk, W. Ford, and D. Solo. *RFC 5280 - Internet X.509 Public Key Infrastructure Certificate and Certificate Revocation List (CRL) Profile*, May 2008.
- [14] G. S. Manku, M. Bawa, and P. Raghavan. Symphony: distributed hashing in a small world. In *USITS*, 2003.
- [15] P. Maymounkov and D. Mazières. Kademlia: A peer-to-peer information system based on the XOR metric. In *IPTPS '02*, 2002.
- [16] J. C. Perez. Facebook's beacon more intrusive than previously thought. http://www.pcworld.com/article/140182/facebook_beacon_more_intrusive_than_previously_thought.html, 2007.
- [17] S. Ratnasamy, P. Francis, S. Shenker, and M. Handley. A scalable content-addressable network. In *Proceedings of ACM SIGCOMM*, 2001.
- [18] S. Ratnasamy, M. Handley, R. M. Karp, and S. Shenker. Application-level multicast using content-addressable networks. In *Workshop on Networked Group Communication (NGC'01)*, 2001.
- [19] A. Rowstron and P. Druschel. Storage management and caching in PAST, a large-scale, persistent peer-to-peer storage utility. In *Symposium on Operating Systems Principles (SOSP'01)*.
- [20] A. Rowstron and P. Druschel. Pastry: Scalable, decentralized object location and routing for large-scale peer-to-peer systems. In *IFIP/ACM International Conference on Distributed Systems Platforms (Middleware)*, November 2001.
- [21] A. Shakimov, H. Lim, L. P. Cox, and R. Caceres. Vis-à-vis: online social networking via virtual individual servers. Technical report, May 2008.
- [22] I. Stoica and et al. Chord: A scalable Peer-To-Peer lookup service for internet applications. In *SIGCOMM*, 2001.
- [23] D. I. Wolinsky, L. Abraham, K. Lee, Y. Liu, J. Xu, P. O. Boykin, and R. Figueiredo. On the design and implementation of structured P2P VPNs. In *TR-ACIS-09-003*, October 2009.
- [24] D. I. Wolinsky, K. Lee, T. W. Choi, P. O. Boykin, and R. Figueiredo. Virtual private overlays: Secure group communication in NAT-constrained environments. In *TR-ACIS-09-004*, December 2009.