

Towards Social Profile Based Overlays

David Isaac Wolinsky, Pierre St. Juste, P. Oscar Boykin, Renato Figueiredo
University of Florida

Abstract—Online social networking has quickly become one of the most common activities of Internet users. As social networks evolve, they encourage users to share more information, requiring the users, in turn, to place more trust into social networks. In centralized systems, this means trusting a third-party commercial entity, like Facebook or MySpace, where as peer-to-peer (P2P) systems can enable the creation of online social networks where trust need only be extended to friends. In this paper, we present a novel approach to constructing completely decentralized social networks through P2P overlays. All users join a common directory overlay, which facilitates friend discovery connecting peers to individualized profile overlays, mapping the ownership of an overlay to the ownership of a profile. Profile management builds upon transparent use of a public key infrastructure (PKI). We define necessary interfaces for constructing the system and describe some examples of user interactions with the system.

I. INTRODUCTION

Online social networking has become pervasive in daily life, though as social networks grow so does the wealth of personal information that they store. Once information has been released on a social network, known as a user's profile, the data and the user are at the mercy of the terms dictated by the social network infrastructure, which today is typically third-party, centrally owned. If the social network engages in activities disagreeable to the user, due to change of terms or opt-out programs not well understood by users such as recent issues with Facebook's Beacon program [1], the options presented to the user are limited: to leave the social network (surrendering their identity and features provided by the social network), to accept the disagreeable activities, or to petition and hope that the social network changes its behavior.

As the use of social networking expands to become the primary way in which users communicate and express their identity amongst their peers, the users become more dependent on the policies of social network infrastructure owners. Recent work [2] explores the coupling between social networks and P2P systems as a means to return ownership to the users, noting that a social network made up of social links is inherently a P2P system with the aside that they are currently developed on top of centralized systems. In this paper, we extend this idea with focus on the topic of topology; that is, how to self-organize social profiles that leverage the benefits offered by a structured P2P overlay abstraction.

Structured P2P overlays provide a scalable, resilient, and self-managing platform for distributed applications. Structured overlays enable users to easily create their own decentralized systems for the purpose of data sharing, interactive activities, and other networking-enabled activities. In this paper, we extend our previous work [3] to enable social network profile

overlays. The previous work addresses the challenges of bootstrapping secure, private overlays in NAT and firewall environments by using a public overlay for discovery and as a relay or communication transport.

Social networks consist of users and groups. Each user has a profile, a set of friends, and private messaging; while each group consists of one or more managers, users, and a messaging board. The profile contains user's personal information, status updates, and public conversations, similar to a message board. Friends are individuals trusted sufficiently by a user to view the user's profile. Private messaging enables sending messages discretely between users without leaking the message to other members. A group is very similar to an individual profile, though it is shared by many users. In some cases, groups are open, though others require approval from an administrator, and in general all users can post messages to the group.

Using this social networking model, we describe how a public overlay can be used as a directory for finding and befriending friends or finding and accessing groups. Once access has been enabled, the public overlay can be used to bootstrap connectivity to existing profile and group overlays. Security for profile are provided by a public key infrastructure (PKI), where profile owners or group managers are the certificate authorities (CA) and all members have signed certificates. The overlay stores profile data or group information in its distributed data store, supporting decentralized access using scalable mechanisms regardless of the profile owner's online presence. In this paper, we present the architecture of these overlays, as presented in Figure 1.




Fig. 1. An example social overlay network. Alice has a friendship with Bob and Carol, hence both are members of her profile overlay. Bob has a friendship with Alice and Dave but not Carol; hence Alice and Dave are members of his profile overlay, while Carol is not. Each peer has many overlay memberships but a single root represented by dashed lines in various shades of gray. For clarity, overlay shortcut connections are not shown.

The rest of this paper is organized as follows. Section II

provides background and related work. Section III describes our multi-overlay approach, explaining how to map social networks onto structured P2P overlays. We express our expectations for user interaction with the system in Section IV. In Section V, we explore some of the remaining challenges introduced by our approach. We conclude the paper in Section VI.

II. BACKGROUND

In this section, we review structured P2P overlays, challenges and solutions for bootstrapping overlays, and other methods for constructing decentralized online social networks. We place an early emphasis on structured P2P systems, as they provide the basis for our approach by creating scalable, autonomic environments, thereby limiting user exposure to the gory details of system configuration and organization. We also focus on the challenges of bootstrapping P2P systems when the system lacks dedicated bootstrapping nodes.

A. Structured P2P Overlays

There are two types of P2P systems, unstructured and structured. Unstructured systems are generally constructed by peers attempting to maintain a certain amount of connections to other peers in the P2P system, whereas structured systems organize into well-defined topologies, such as 2-D rings or hypercubes. Though unstructured systems are typically simpler to bootstrap and maintain, they rely on global knowledge, flooding, or stochastic techniques to search for information in an overlay and thus have scalability constraints. Alternatively, structured systems have guaranteed search time typically with a lower bound of $O(\log N)$, though recent research has shown that these systems can be constructed to support $O(1)$ search time. Some examples of structured systems are Pastry [4], Chord [5], Symphony [6], Kademlia [7], CAN [8], and Dynamo [9].

A key component of most structured overlays is support for decentralized storage / retrieval of information by mapping keys to specific node IDs in an overlay called a distributed hash table (DHT). At a minimum, the data is stored at the node ID either smaller or larger to the data's node ID and for fault tolerance the data can be stored at other nodes. DHTs can be used by peers of systems to coordinate allocation and discovery of resources, making them attractive for self-configuration in decentralized collaborative environments. DHTs provide the building blocks to form more complex distributed data stores as presented in Past [10] and Kosha [11].

B. Bootstrapping P2P Overlays

One of the key challenges to overlays is the bootstrapping problem, that is, how does one find an active member in the overlay in order to become a member. In large public systems, often times, the owners of the project and sometimes corporate sponsors provide these resources, though these types of bootstrapping peers are typically not available for small, private overlays. The bootstrapping problem can be divided into two components: finding a bootstrapping member of the

overlay and handling connectivity constraints if the overlay has no members with public addresses. In [12, 13], the authors discuss the concept of a single overlay supporting services through the use of additional overlays, which use the underlying overlay to assist in discovery, this deals with the component of finding a bootstrap peer. To address the challenge of bootstrapping peers being behind NATs, we described a system in [3] enables the public overlay to act as a NAT traversal rendezvous point. More specifically, a peer, attempting to join a private overlay, queries the public overlay's DHT to obtain a list of currently online private overlay peers who have matching public IDs. The peers then use the public overlay, in addition to direct UDP or TCP links, to communicate with the private overlay nodes and bootstrap into the private overlay. This method enables peers even behind firewalls and NATs to have publicly available addresses in the public overlay. Additionally this work describes mechanisms that enable the use of PKI technologies to seamlessly secure both point-to-point and end-to-end communication in a structured overlay.

During evaluation, using both simulated and real systems, the time for a single peer to join first the public and thereafter private overlays was small and grew logarithmically with network sizes. The real system was tested using a public overlay of 600 nodes on PlanetLab and a random distribution of peers in the private overlay. With point-to-point security links enabled in the private overlay, the time to connect was less than 22 seconds for all cases. In simulations, overlays with as many as 100,000 peers were evaluated. For the 100,000-peer overlay, regardless of the private overlay size and with security enabled, peers were able to connect to their private overlays in less than 48 seconds. In relation to this paper, these results can be interpreted such that the latency required for a single peer to from being completely disconnected from the social network to being fully connected to the directory and all profile overlays.

C. Peer-to-Peer Social Networks

In [14], a DHT provides the look-up service for storing meta data pertaining to a peer's profile. Peers query the DHT for updated content from their friends by hashing their unique identifiers (e.g. friends' email addresses). The retrieved meta data contains information for obtaining the profile data such as IP address and file version. Their work relies on a PKI system that provides identification, encryption, and access control. In contrast, our approach provides each user their own private overlay secured by point-to-point encryption and authentication amongst all peers in the profile overlay. The profile overlay provides a clean abstraction of access control, whereby once admitted to a private overlay, users can access a distributed data store which holds the contents of the owners profile.

[15] takes a different approach by depending on virtual individual servers (VIS) hosted on a cloud infrastructure such as Amazon EC2. Friends contact each other's VIS directly for updates. A DHT is used as a directory for groups and

interest-based searches. Their approach assumes bidirectional end-to-end connectivity between each VIS, where a profile is only available during the up time of the VIS. Because of the demands on network connectivity and up time, the approach assumes a cloud-hosted VIS and has difficulty being used on user-owned resources. Our approach enables users to avoid the need for all-to-all connectivity and constant up time through the use of NAT traversal support and the ability to store the profile in the overlay's distributed data store.

The approach presented in [16] relies on a central system to host identities and certificates that can then be used to query a DHT to discover an initial hop in a route to a specific peer through their circle of friends. The circle of friends consists of an unstructured overlay, where direct friends maintain direct connections with the peer, and outer circles consist of friends of friends and friends of friends of friends. The main goal of this work is to remove the private components of a profile from a central entity, whereas our approach makes a clean break from all centralization and emphasizes scalability through distributed replica techniques.

Unlike the above approaches, the P2P social network presented in [17] uses an unstructured overlay without a DHT where peers connect directly to each other rather than through the overlay establishing unique identifiers to deal with dynamic IPs. Peers cache each other's data to improve availability. While helper nodes are used to assist with communication between peers behind NATs. The approach lacks security and access control considerations and lacks the guarantees and the simplicity of the abstraction offered by a structured overlay.

III. SOCIAL OVERLAYS

In this section, we explain how to map online social networking to virtual private overlay based social network consisting of a public directory overlay with many private profile overlays. The directory overlay supports friend discovery and verification and stores a lists of peers currently active in each profile overlay. Profile overlays support message boards, private messages, and media sharing.

A. Finding and Verifying Friends

In a traditional social network, a directory provides the ability to search for users using public information, such as the user's full name, user ID, e-mail address, group affiliations, and friends. The search results return zero or more matching directory entries. Based upon the results, the user, *A*, can potentially make a friendship request. The request receiver, *B*, can review the public information of *A* to making a decision. If *B* accepts the request, *A* and *B* are given access to each other's profiles. Once profile access has been enabled, the users can learn more information, and if it turns out to be a mistake, the peers can unilaterally end the relationship.

One method to map this to our proposed social overlay would be for directory entries to be inserted into the DHT of a public overlay. A user could store their public information at a single location in the DHT, which would be indexed through multiple mappings involving their public information.

For example, a user could store at the DHT key $hash("alice")$ or $hash("alicebob")$ a pointer to the DHT location for their certificate. The key here is that any subset of the user's public information in lower-case format could be hashed into an index into the DHT that would eventually direct the searching user to one or more user's certificates. Public information should consist of as much identifiable information the peer is willing to share and an overlay address. The overlay address enables asynchronous offline messaging, such as friendship requests, which will be explained shortly.

Because the users need a way to verify each other that involves social credentials, we propose the use of a new form of certificate. The main portion of the certificate is similar to a self-signed x509 certificate with public information such as user's name, e-mail address, and group affiliations embedded into the certificate. At the tail of the certificate is a list of friend authorizations using to techniques similar to PGP. The field should consist of a hash of the user's certificate, a hash of the friend's certificate, and a time stamp all signed by the friend's certificate. Since PGP does not provide a strong method for revocation, the time stamp assists in deciding whether or not a friendship link is still active without accessing the profile overlay of either peers. Thus peers should actively maintain these lists, so that no entries appear to be invalid.

While looking for an individual, a peer may discover that many individuals have overlapping public information components, such as the user's name. Assuming all entries are legitimate, the overlay must have some method of supporting multiple, distinct values at the same key, leaving the peer or the peer's DHT client to parse the responses and determining the best match by reviewing the contents of each certificate. Alternatively, a technique like Sword [18], which supports distributing the data across a set of nodes and using a bounded broadcast to discover peers that match all information, could be used for searching.

If a peer, *A*, desires a friendship with another peer, *B*, *A* issues a friendship request, which will be stored in the DHT at the overlay address listed in *B*'s certificate, as described earlier. The friendship request consists of the self-signed certificate of *A*, the requesting peer; the public information of the request receiver, *B*; and a time stamp; all signed with the private key associated with *A*'s private key matched to their self-signed certificate.

Within a reasonable amount of time after a request has been inserted into the DHT, *B* can come online and check for outstanding requests. Upon receiving a request, *B* has three choices: a conditional accept, an unconditional accept, or a reject. During an unconditional accept, *B* signs *A*'s request and issues a request to befriend *A*. Alternatively in the case of a conditional accept, *B* issues a friendship request, waits for a reply, and investigates the profile prior to signing the *A*'s request. Once a user has received a signed certificate, they may access the remote peer's profile overlay as discussed in III-B, which is also responsible for activities such as revocation.

Discovery of a user is not limited to the directory entries. Because users have a public overlay based mailbox that

does not easily identify them, the system can act as an anonymous message relay. This enables mechanisms out of band discovery, key for users that do not feel comfortable with placing information in the directory overlay. Peers can share their mailbox identification and certificates through mechanisms like e-mail, chat, or personal websites. Once a peer has a mailbox identification and , they can submit secure friendship requests.

B. The Profile Overlay

In a traditional social network, the profile or user-centric portion consists of private messaging, data sharing, friendship maintenance, and a public message board for status updates or public messages. In this section, we explain how these components can be applied to a structured overlay dedicated to an individual profile.

Using the techniques such as those described in [3], it is feasible to efficiently multiplex a P2P system across multiple, virtual private overlays enabling each profile owner to have a profile overlay consisting of their online friends. For access control, we employ a PKI, where each member uses the signed certificate generated during the “finding and verifying friends” stage. All links are encrypted using symmetric security algorithms established through the PKI, thus preventing uninvited guests from gaining direct access to the overlay and hence the profile. Because the profile owner also is the CA for all members of the overlay, they can easily revoke users from access to the profile overlay. [3] describes efficient mechanisms for overlay revocation through the use of broadcasting for immediate revocation and the use of DHT for indirect and permanent revocation.

The message board of a profile can be stored in two ways: distributed within the profile overlay via a data store or stored on the profile owner’s personal computing devices. The distributed data store provide the profile when the owner is offline and also distributes the load for popular profiles. For higher availability, each peer should always be a provider for all data in their profile when they are online. To ensure authenticity and integrity, all peers should sign their messages and each peer’s certificate should be available in the overlay for verification. Messages that are unsigned should be ignored by all members of the overlay. An ideal overlay for this purpose should support complex queries [19] allowing easy access to data stored chronologically, by content, by type, i.e., media, status updates, or message board discussions.

Private messaging in the profile overlay is unidirectional; only the profile owner can receive private messages using their overlay. To enforce this, a private message should be prepended with a symmetric key encrypted by the profile owners public key, the message should be appended by a hash of the message to ensure integrity and the entire message encrypted by the symmetric key. This approach ensures that only the sender and the profile owner can decrypt the private message. The contents of the private message should include the sender, time sent, and the subject. Messages can be stored in well known locations, so that the profile owner can either

poll the location or, alternatively, use an event based system to notify them of the new message.

C. Event Based Message Notification

Both the directory and profile overlays have methods by which peers can receive messages. In the directory overlay, these take form by means of friendship request and accept. In the profile overlay, they consist of private messages. While polling the location in the DHT occasionally will allow peers to receive the messages, polling has inherent delays and network costs. Alternatively, an event mechanism would enable peers to receive sent messages very quickly after they have been sent with minimal impact on network throughput.

A simple method for implementing an event notification system involves using the DHT. Each event would have an identification that would map to a list of peers wanting to know when an event occurred and the data associated with it. Thus mapping the *(eventid, listener)* to the DHT could be done by hashing a string such as “private messages for me” and storing the profile owners active nodes into the list of listeners. When a message was stored to the users mailbox, the sender could query this list and send to each listener a notification of the new private message. Alternatively, if a higher degree of anonymity is required, the DHT server could be modified to forward the response to the listeners directly rather than returning a list of listeners. Of course, this does not prevent potential race conditions occurring, such as a situation where a peer recently joined their profile overlay, had already queried their mailbox and found it empty, while simultaneously a private message was sent to them yet they were not in the listeners list. Thus occasional polling is required, though can be minimized, the longer a node has been online.

D. Active Peers

The directory overlay should be used to assist in finding currently active peers in the profile overlays. By placing their node IDs at a well-known, unique per-profile overlay keys in the DHT, active peers can bootstrap incoming peers into the profile overlay. We implemented and evaluated this concept in [3]. Because the profile overlay members all use PKI to ensure membership, even if malicious peers insert their ID into the active list, it would be useless as the peer would only form connections with peers who also have a signed certificate.

E. Groups

Groups can be considered extensions of profile overlays. The fundamental difference between a group and a profile is that a group lacks private messaging and has shared ownership. So just as a peer can find a profile in the directory by hashing the name of the user and other identifiable information, so can the user find the group. Like the certificate of the user, the members of a group sign the groups certificate to represent their membership to that group. The user requests membership to the group and a group manager can sign the certificate allowing that member access to the group. Finally, the group

can be bootstrapped in the same way as the profile overlay through the directory overlay.

The unique challenge presented by groups is the sharing of the CA task. A decentralized solution would be for all members of the group to be listed in the groups DHT and when a peer becomes a manager, their certificate is added to the certificate chain for the group enabling them to sign certificates for that group. If an administrator loses their position, then all members who had their certificate signed by that administrator would need to obtain a new certificate. Alternatively, the owner of the group and the original CA could sign the group members certificates again, so that they are signed by the root CA and would not lose access if administration changes.

IV. USER INTERACTION

Though this paper describes a complex organization of a social network, the result should be transparent to the user such that interaction with this system should be no different than existing online social networks. The user space implementation could consist of a downloadable application or a browser based Flash or Silverlight application. If the user, Bob, had already created an account, Bob would be presented with an interface showing their friends profiles. Based upon Bob's configuration, the social application could retrieve profile updates as he navigates to individual profiles or as soon as the application joins an individual profile overlay, reactive versus proactive profile querying.

If this was Bob's first time starting the application, he would be presented with screens asking for his privacy preferences, such as whether or not he wants his information in the directory overlay, if he felt comfortable enough with the idea of people knowing he was a member of the social network and who his friends are. Then application would ask for personal information to populate his profile and to generate his directory information. At which point, the application would join the overlay and create Bob's private overlay. Bob could then start searching for friends, make friend requests, and respond to friend requests.

Recently, Bob had been thinking about his high school days and was curious if Alice was also a member of this social network, though Bob did not have Alice's e-mail address, just her first and last name. Bob enters Alice's name into the search box and is presented by a list of Alice's. As Bob reviews each of the entries, he recognizes an Alice that is friend's with some of the same people Bob was in high school. Bob selects to become her friend. At which point, the application transparently inserts a friendship request to Alice and signs Alice's certificate so Alice can view Bob's profile. Of course that is because Bob has chosen to allow user-initiated friend requests access to his profile. Alice receives Bob's request, peruses his profile and feels fine becoming friends with Bob, which initiates a transparent process of signing Bob's certificate and placing the result in the public overlay. There is one problem though, when Bob receives Alice's signature and views her profile, he realizes that this is some other Alice. He quickly chooses to defriend her. This

causes Bob's application to broadcast a revocation for Alice's signature and to store the revocation in the DHT. Alice, who was viewing Bob's profile, is notified of this sudden loss of trust and while she is able to view the contents of Bob's profile, which she has already accessed and obtained, she can no longer receive updates as members of Bob's overlay prevent her from accessing it.

In another instance, Bob bumped into Carol, who e-mailed Bob a copy of her profile blob, Carol's public overlay mailbox and certificate. Bob points his social application to this configuration data, and it immediately submits a request to become Carol's friend. Carol receives notification and accepts Bob's friendship request. At this point, both Bob and Carol have transparently exchanged signed certificates and have mutual access to each other profiles. As Bob reads Carol's latest news, he remembers a funny personal story and that he would like to privately share with Carol. So he sends Carol a private message. Carol is offline though. The next time Carol goes online, her social application discovers the message and presents it to her. In this scenario, Bob's application has taken the private message, secured it with her public key and a symmetric key. After which, it inserts it into the DHT and sends a notice to the event notification system, which detects that there were no entries listening for incoming messages. When Carol's application comes online, it queries the DHT receiving the message. The message is decrypted and authenticated ensuring that it originated from Bob and then presents the results to Carol.

V. CHALLENGES

While structured P2P overlays have been well-studied in a variety of applications, their use in social profile overlays raises new interesting questions, including:

1) Handling small overlay networks - P2P overlay research typically focuses on networks larger than the typical user's friend count (Facebook's average is 130¹). Because social profile overlays are comparatively smaller, this can impact the reliability of the overlay and availability of profile data. A user can host their own profile; however when the user is disconnected it is important that their profile remains available even under churn. It is thus important to characterize churn in this application to understand how to best approach this problem. An optional of per-user deployment of a virtual individual server (VIS) and the use of replication schemes aware of a user's resources provide possible directions to address this issue.

2) Overlay support for low throughput, unconnected devices - devices such as smart phones cannot constantly be actively connected to the overlay and the connection time necessary to retrieve something like a phone number may be too much to make this approach useful. Similar to the previous challenge, this approach could benefit from using a VIS enabling users access to their social overlays by proxy without establishing a direct connection to the overlay network.

¹<http://www.facebook.com/press/info.php?statistics>

3) Reliability of the directory and profile overlay -

Overlays are susceptible to attacks that can nullify their usefulness. While the profile overlay does have point-to-point security, in the public, directory overlay, the lack of any form centralization makes policing the system a complicated procedure. While our approach of appending friends list can assist users in making decisions on identity, it does not protect against denial of service attacks. For example, users could attempt create many similar identities in an attempt to overwhelm a user in their attempt to find a specific peer. Previous work has proposed methods to ensure the usability of overlays even while under attack. For the social overlay to be successful, we must identify which methods should be used. A possible approach is to replicate public information within a user's profile overlay thus providing an alternative directory overlay for querying prior to using the public directory overlay.

VI. CONCLUSION

In this paper, we proposed methods by which a social network can be decentralized through the use of structure P2P overlays. P2P systems, in general, have very nice properties that make them attractive for average users. Namely, they require minimal configuration and organization of resources, in addition, P2P systems naturally handle increased demand to additional peers. Our approach is based upon the use of a multiple overlay system, where all users join a public directory overlay and their's as well as their friends profile overlays. The directory overlay enables users to find and befriend other peers and bootstrap connections into the secure profile overlays. Upon forming a friendship through the directory overlay, peers are given CA signed certificates that allow them to join each other's profile overlay. The owner of the profile overlay acts as CA enabling unilateral dismissal of friendships via certificate revocation using efficient and reliable methods. Additionally, the profile overlay provides a clean abstraction for privacy purposes that simplifies access to profile data, as only members of the overlay have access to the profile. For the purpose of storing profile information into the overlay, we cite previous work that can be used to provide distributed data services and give examples of how to store data securely in the overlay. Our proposed system returns control of the social network and more importantly users' identity to the users and eliminates the need for centralized social networks.

REFERENCES

- [1] J. C. Perez, "Facebook's beacon more intrusive than previously thought," http://www.pcworld.com/article/140182/facebook_beacon_more_intrusive_than_previously_thought.html, 2007.
- [2] S. Buchegger and A. Datta, "A case for P2P infrastructure for social networks - opportunities & challenges," in *WONS '09: The Sixth International Conference on Wireless On-demand Network Systems and Services*, 2009.
- [3] D. I. Wolinsky, K. Lee, T. W. Choi, P. O. Boykin, and R. Figueiredo, "Virtual private overlays: Secure group communication in NAT-constrained environments," January 2010.
- [4] A. Rowstron and P. Druschel, "Pastry: Scalable, decentralized object location and routing for large-scale peer-to-peer systems,"

- in *IFIP/ACM International Conference on Distributed Systems Platforms (Middleware)*, November 2001.
- [5] I. Stoica and et al., "Chord: A scalable Peer-To-Peer lookup service for internet applications," in *SIGCOMM*, 2001.
- [6] G. S. Manku, M. Bawa, and P. Raghavan, "Symphony: distributed hashing in a small world," in *USITS*, 2003.
- [7] P. Maymounkov and D. Mazières, "Kademlia: A peer-to-peer information system based on the XOR metric," in *IPTPS '02*, 2002.
- [8] S. Ratnasamy, P. Francis, S. Shenker, and M. Handley, "A scalable content-addressable network," in *In Proceedings of ACM SIGCOMM*, 2001.
- [9] G. DeCandia, D. Hastorun, M. Jampani, G. Kakulapati, A. Lakshman, A. Pilchin, S. Sivasubramanian, P. Voshall, and W. Vogels, "Dynamo: amazon's highly available key-value store," in *SOSP '07: Proceedings of twenty-first ACM SIGOPS symposium on Operating systems principles*. New York, NY, USA: ACM, 2007, pp. 205–220.
- [10] A. Rowstron and P. Druschel, "Storage management and caching in PAST, a large-scale, persistent peer-to-peer storage utility," in *Symposium on Operating Systems Principles (SOSP'01)*.
- [11] A. R. Butt, T. A. Johnson, Y. Zheng, and Y. C. Hu, "Kosha: A peer-to-peer enhancement for the network file system," in *IEEE/ACM Supercomputing 2004*.
- [12] M. Castro, P. Druschel, A.-M. Kermarrec, and A. Rowstron, "One ring to rule them all: Service discover and binding in structured peer-to-peer overlay networks," in *SIGOPS European Workshop*, Sep. 2002.
- [13] S. Ratnasamy, M. Handley, R. M. Karp, and S. Shenker, "Application-level multicast using content-addressable networks," in *Workshop on Networked Group Communication (NGC'01)*.
- [14] S. Buchegger, D. Schiöberg, L. H. Vu, and A. Datta, "Peerson: P2p social networking: early experiences and insights," in *SNS '09: Proceedings of the Second ACM EuroSys Workshop on Social Network Systems*, 2009.
- [15] A. Shakimov, H. Lim, L. P. Cox, and R. Caceres, "Vis-à-vis:online social networking via virtual individual servers," Tech. Rep., May 2008.
- [16] L. A. Cuttillo, R. Molva, and T. Strufe, "Privacy preserving social networking through decentralization," in *Wireless On-Demand Network Systems and Services (WONS'09)*.
- [17] S. M. A. Abbas, J. A. Pouwelse, D. H. J. Epema, and H. J. Sips, "A gossip-based distributed social networking system," in *Enabling Technologies, IEEE International Workshops on*, 2009.
- [18] J. Albrecht, D. Oppenheimer, A. Vahdat, and D. A. Patterson, "Design and implementation trade-offs for wide-area resource discovery," in *ACM Trans. Internet Technol.*, 2008.
- [19] M. Harren, J. M. Hellerstein, R. Huebsch, B. T. Loo, S. Shenker, and I. Stoica, "Complex queries in dht-based peer-to-peer networks," in *IPTPS '01: Revised Papers from the First International Workshop on Peer-to-Peer Systems*, 2002.